

Advanced Operating System with UNIX

Unit 3 Processes

Christian Kniep

Internation Center of Applied Technologies Bandung

6. August 2010

Table of content

- 1 3.1 Introduction
 - address space & data structure
 - what defines a process
- 2 3.2 Processes Priority
 - Scheduling
- 3 3.3 Other Related Commands
- 4 3.4 The fork() System Call
 - Explanation
- 5 3.5 Child Process Termination
 - relationships
- 6 3.6 Inter Process Communication
 - 3.6.1 Pipes
 - 3.6.2 Named Pipes
 - 3.6.2 Message Queues
 - 3.6.3 Semaphores
 - 3.6.4 Shared Memory

What is a process?

- A process under UNIX consists of an address space and a set of data structures in the kernel to keep track of that process.

What is a process?

- A process under UNIX consists of an address space and a set of data structures in the kernel to keep track of that process.
- The memory-section contains:
 - ▶ the code to execute
 - ▶ the process stack

What is a process?

- A process under UNIX consists of an address space and a set of data structures in the kernel to keep track of that process.
- The memory-section contains:
 - ▶ the code to execute
 - ▶ the process stack
- The kernel keeps track of:
 - ▶ address space map
 - ▶ current status of the process
 - ▶ execution priority of the process
 - ▶ resource usage of the process
 - ▶ current signal mask
 - ▶ owner of the process

what defines a process

what you might see

- A process has certain attributes, including these:
 - PID The unique identifier of a process

what defines a process

what you might see

- A process has certain attributes, including these:
 - PID The unique identifier of a process
 - PPID The parents PID (from whom the process are started)

what defines a process

what you might see

- A process has certain attributes, including these:
 - PID** The unique identifier of a process
 - PPID** The parents PID (from whom the process are started)
 - UID** the UserID of the user who owns the process

what defines a process

what you might see

- A process has certain attributes, including these:
 - PID The unique identifier of a process
 - PPID The parents PID (from whom the process are started)
 - UID the UserID of the user who owns the process
 - GID the GroupID of the owner

what defines a process

what you might see

- A process has certain attributes, including these:
 - PID The unique identifier of a process
 - PPID The parents PID (from whom the process are started)
 - UID the UserID of the user who owns the process
 - GID the GroupID of the owner
 - EUID the EffectiveUserID
 - EGID the EffectiveGroupID

what defines a process

what you might see

- A process has certain attributes, including these:
 - PID The unique identifier of a process
 - PPID The parents PID (from whom the process are started)
 - UID the UserID of the user who owns the process
 - GID the GroupID of the owner
 - EUID the EffectiveUserID
 - EGID the EffectiveGroupID
 - ▶ WTF, what is that suppose to mean!?

what defines a process

what you might see

- A process has certain attributes, including these:

PID The unique identifier of a process

PPID The parents PID (from whom the process are started)

UID the UserID of the user who owns the process

GID the GroupID of the owner

EUID the EffectiveUserID

EGID the EffectiveGroupID

- ▶ WTF, what is that suppose to mean!?

Priority Priority (urgency) the process runs at

what defines a process

The ps-Command

- The command `ps -l`

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
4	R	0	6843	6842	0	80	0	-	1065	-	pts/0	00:00:00	bash
0	R	0	7552	6843	0	80	0	-	872	-	pts/0	00:00:00	ps

what defines a process

The ps-Command

- The command `ps -l`

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
4	R	0	6843	6842	0	80	0	-	1065	-	pts/0	00:00:00	bash
0	R	0	7552	6843	0	80	0	-	872	-	pts/0	00:00:00	ps

- F** (PROCESS FLAGS)
 - 1** forked but didn't exec
 - 4** used super-user privileges
 - 30** according to the book 'loaded into memory'

what defines a process

The ps-Command

- The command `ps -l`

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
4	R	0	6843	6842	0	80	0	-	1065	-	pts/0	00:00:00	bash
0	R	0	7552	6843	0	80	0	-	872	-	pts/0	00:00:00	ps

- F** (PROCESS FLAGS)

- 1** forked but didn't exec
- 4** used super-user privileges
- 30** according to the book 'loaded into memory'

- S** (PROCESS STATE CODES)

- D** Uninterruptible sleep (usually IO)
- R** Running or runnable (on run queue)
- S** Interruptible sleep (waiting for an event to complete)
- Z** Defunct (Bombie") process, terminated but not reaped by its parent

what defines a process

The ps-Command

● **PID** The ProcessID

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRP)** The ProcessGRouPID)

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRP)** The ProcessGRouPID)
- **PRI** The priority of the process

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRPID** The ProcessGRouPID)
- **PRI** The priority of the process
- **NI** Nice value; used in priority computation
 - ▶ the lower the value of PRI and NI the higher the priority

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRP)** The ProcessGRouPID)
- **PRI** The priority of the process
- **NI** Nice value; used in priority computation
 - ▶ the lower the value of PRI and NI the higher the priority
 - ▶ if a process uses the CPU the PRI-value will raise the nice value

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRPID** The ProcessGRouPID)
- **PRI** The priority of the process
- **NI** Nice value; used in priority computation
 - ▶ the lower the value of PRI and NI the higher the priority
 - ▶ if a process uses the CPU the PRI-value will raise the nice value
- **P** Which processor runs the process

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRP)** The ProcessGRouPID)
- **PRI** The priority of the process
- **NI** Nice value; used in priority computation
 - ▶ the lower the value of PRI and NI the higher the priority
 - ▶ if a process uses the CPU the PRI-value will raise the nice value
- **P** Which processor runs the process
- **SZ (VSZ)** (Virtual) Memory is used

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRP)** The ProcessGRouPID)
- **PRI** The priority of the process
- **NI** Nice value; used in priority computation
 - ▶ the lower the value of PRI and NI the higher the priority
 - ▶ if a process uses the CPU the PRI-value will raise the nice value
- **P** Which processor runs the process
- **SZ (VSZ)** (Virtual) Memory is used
- **TTY** The terminal the process is running at

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRP)** The ProcessGRouPID)
- **PRI** The priority of the process
- **NI** Nice value; used in priority computation
 - ▶ the lower the value of PRI and NI the higher the priority
 - ▶ if a process uses the CPU the PRI-value will raise the nice value
- **P** Which processor runs the process
- **SZ (VSZ)** (Virtual) Memory is used
- **TTY** The terminal the process is running at
- **Time** The cumultative execution time

what defines a process

The ps-Command

- **PID** The ProcessID
- **PPID** The ParentProcessID
- **(PGRP)** The ProcessGRouPID)
- **PRI** The priority of the process
- **NI** Nice value; used in priority computation
 - ▶ the lower the value of PRI and NI the higher the priority
 - ▶ if a process uses the CPU the PRI-value will raise the nice value
- **P** Which processor runs the process
- **SZ (VSZ)** (Virtual) Memory is used
- **TTY** The terminal the process is running at
- **Time** The cumultative execution time
- **COMD** The executed command

resources

- reward patience
 - ▶ those who have used the least will get access served

resources

- reward patience
 - ▶ those who have used the least will get access served
 - ▶ those who waits for an event will get higher prio (e.g. keyboard)

resources

- reward patience
 - ▶ those who have used the least will get access served
 - ▶ those who waits for an event will get higher prio (e.g. keyboard)
- The CPU is the most wanted resource. Tied together with the other (such as RAM, Harddrive, Network, etc.) it defines the throughput of an machine

resources

- reward patience
 - ▶ those who have used the least will get access served
 - ▶ those who waits for an event will get higher prio (e.g. keyboard)
- The CPU is the most wanted resource. Tied together with the other (such as RAM, Harddrive, Network, etc.) it defines the throughput of an machine
- nice

resources

- reward patience
 - ▶ those who have used the least will get access served
 - ▶ those who waits for an event will get higher prio (e.g. keyboard)
- The CPU is the most wanted resource. Tied together with the other (such as RAM, Harddrive, Network, etc.) it defines the throughput of an machine
- nice
 - ▶ the nice-command allows the user to manipulate the internal scheduler

resources

- reward patience
 - ▶ those who have used the least will get access served
 - ▶ those who waits for an event will get higher prio (e.g. keyboard)
- The CPU is the most wanted resource. Tied together with the other (such as RAM, Harddrive, Network, etc.) it defines the throughput of an machine
- nice
 - ▶ the nice-command allows the user to manipulate the internal scheduler
 - ▶ only a administrator can set a higher priority (lower value)

resources

- reward patience
 - ▶ those who have used the least will get access served
 - ▶ those who waits for an event will get higher prio (e.g. keyboard)
- The CPU is the most wanted resource. Tied together with the other (such as RAM, Harddrive, Network, etc.) it defines the throughput of an machine
- nice
 - ▶ the nice-command allows the user to manipulate the internal scheduler
 - ▶ only a administrator can set a higher priority (lower value)
 - ▶ depending on the Unix-flavor the the increments and syntax is different

resources

`top` have a live updating view on the system processes

resources

`top` have a live updating view on the system processes

- ▶ There are many childs of `top` with different approaches

`htop` more information with better formatting

`ntop` network-performance (what process does what)

... and so on

resources

top have a live updating view on the system processes

- ▶ There are many childs of top with different approaches

htop more information with better formatting

ntop network-performance (what process does what)

... and so on

lsof to have a look what files are opened (maybe have to be installed first)

resources

top have a live updating view on the system processes

- ▶ There are many childs of top with different approaches

htop more information with better formating

ntop network-performance (what process does what)

... and so on

lsof to have a look what files are opened (maybe have to be installed first)

fuser what process is using certain files

What is it?

- `fork()` creates a new process

What is it?

- `fork()` creates a new process
- the memory allocated to the process will be duplicated

What is it?

- `fork()` creates a new process
- the memory allocated to the process will be duplicated
- from this line on both processes running the same code

Who was first?

- if a child dies, the parent has to recognize it

Who was first?

- if a child dies, the parent has to recognize it
- if the parent don't, the child will becomm a 'zombie'

Whats out there?

- **Pipes** Point to point, only one way

Whats out there?

- **Pipes** Point to point, only one way
- **Named Pipes** Looks like a file, thing about it like a spot

Whats out there?

- **Pipes** Point to point, only one way
- **Named Pipes** Looks like a file, thing about it like a spot
- **Message queues** FIFO

Whats out there?

- **Pipes** Point to point, only one way
- **Named Pipes** Looks like a file, thing about it like a spot
- **Message queues** FIFO
- **Semaphores** Guarded line

Whats out there?

- **Pipes** Point to point, only one way
- **Named Pipes** Looks like a file, thing about it like a spot
- **Message queues** FIFO
- **Semaphores** Guarded line
- **Shared Memory** Use the same addresses

Whats out there?

- **Pipes** Point to point, only one way
- **Named Pipes** Looks like a file, thing about it like a spot
- **Message queues** FIFO
- **Semaphores** Guarded line
- **Shared Memory** Use the same addresses
- **Sockets** FIFO in a network

3.6.1 Pipes

Point to Point with FIFO

- One-Way communication from one entity to another

3.6.1 Pipes

Point to Point with FIFO

- One-Way communication from one entity to another
- Has an INode but no reference on it, so its **NOT** a file

3.6.1 Pipes

Point to Point with FIFO

- One-Way communication from one entity to another
- Has an INode but no reference on it, so its **NOT** a file
- with `read` and `write` you could use it

A Pipe With a Name

- Like pipes, but both ways and

A Pipe With a Name

- Like pipes, but both ways and
- they are a file

3.6.2 Named Pipes

A Pipe With a Name

- Like pipes, but both ways and
 - they are a file
- ⇒ You have to `open()` it first

a bit organized

a bit organized

- messages have types (number,msg)
- you could read only type-number X

a bit organized

- messages have types (number,msg)
- you could read only type-number X
- could be private (only parent and childs) or public

3.6.3 Semaphores

guard

- its like a queue with a counter

3.6.3 Semaphores

guard

- its like a queue with a counter
- it could seperate enemies

3.6.3 Semaphores

guard

- its like a queue with a counter
- it could seperate enemies
- or unites friends

3.6.3 Semaphores

guard

- its like a queue with a counter
- it could seperate enemies
- or unites friends
- its based on a flag and a queue

on the edge

- if you want be fast, why transfer it?

on the edge

- if you want be fast, why transfer it?
- if you use the same memory-address you don't have to

over the network

- like named pipes not based on files, but on network-ports